

# Fs

## Design

### 0x1 doubly-indirect block

Modify `bmap()` so that it implements a doubly-indirect block, in addition to direct blocks and a singly-indirect block. You'll have to have only 11 direct blocks, rather than 12, to make room for your new doubly-indirect block; you're not allowed to change the size of an on-disk inode. The first 11 elements of `ip->addrs[]` should be direct blocks; the 12th should be a singly-indirect block (just like the current one); the 13th should be your new doubly-indirect block. You are done with this exercise when `bigfile` writes 65803 blocks and `usertests` runs successfully:

Xv6 file system uses ext-like 12 direct blocks and 1 singly-indirect block in its inodes.

my job is to add a doubly-indirect block index in xv6 inode

1. modify inode structure:

```
1  #define NDIRECT 11
2  #define NINDIRECT (BSIZE / sizeof(uint))
3  #define MAXFILE (NDIRECT + NINDIRECT + NINDIRECT*NINDIRECT)
4
5  // On-disk inode structure
6  struct dinode {
7      short type;           // File type
8      short major;          // Major device number (T_DEVICE only)
9      short minor;          // Minor device number (T_DEVICE only)
10     short nlink;           // Number of links to inode in file system
11     uint size;             // Size of file (bytes)
12     uint addrs[NDIRECT+2]; // Data block addresses
13 };
```

2. modify `bmap()`

`bmap()` return the disk block address of the nth block in inode ip. my job is to add the search in doubly-indirect index.

```
1  static uint
2  bmap(struct inode *ip, uint bn)
3  {
4      //printf("handle %d\n",bn);
5      uint addr, *a;
```

```

6  struct buf *bp;
7  //if the data block in direct block
8  ...
9  //if data block in indirect block
10 ...
11 //if data block in double indirect block
12 bn -= NINDIRECT;
13 if(bn < NINDIRECT * NINDIRECT) {
14     //printf("alloc in double block\n");
15     //locate double indirect block
16     uint l1 = bn / NINDIRECT;
17     uint l2 = bn % NINDIRECT;
18     //load l0 double-indirect block, allocating if necessary
19     if((addr = ip->addrs[NDIRECT+1]) == 0)
20         ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
21     bp = bread(ip->dev, addr);
22     a = (uint*)bp->data;
23     uint addr1;
24     if ((addr1 = a[l1]) == 0){
25         a[l1] = addr1 = balloc(ip->dev); // save l1 indirect block in addr1
26         log_write(bp);
27     }
28     brelse(bp);
29     //load l1 indirect block, allocating if necessary
30     bp = bread(ip->dev, addr1);
31     a = (uint*)bp->data;
32     if((addr = a[l2]) == 0){
33         a[l2] = addr = balloc(ip->dev);
34         log_write(bp);
35     }
36     brelse(bp);
37     return addr;
38 }
39 panic("bmap: out of range");
40 }

```

## 0x2 symbolic link

You will implement the `symlink(char *target, char *path)` system call, which creates a new symbolic link at `path` that refers to file named by `target`. For further information, see the man page `symlink`. To test, add `symlinktest` to the Makefile and run it. Your solution is complete when the tests produce the following output (including `usertests` succeeding).

1. write a new type "symlink" record in inode

```

1  uint64 sys_symlink(void){
2

```

```

3  char name[DIRSIZ], new[MAXPATH], old[MAXPATH];
4  struct inode *dp, *ip;
5
6  if(argstr(0, old, MAXPATH) < 0 || argstr(1, new, MAXPATH) < 0)
7      return -1;
8  begin_op();
9  //printf("from %s to %s\n", old,new);
10 ip = namei(old);
11 dp = namei(new);
12 if ( ip !=0 && ip->type != T_DIR ){
13     // add a link to the old path
14     ilock(ip);
15     ip->nlink++;
16     iupdate(ip);
17     iunlockput(ip);
18 }
19 if (dp==0){
20     //fsprintf("dp is null\n");
21     if (nameiparent(new,name)==0){
22         printf("a null path under a null parent\n");
23         goto bad;
24     }
25     if((dp= create(new,T_SYMLINK,0,0)) == 0){
26         printf("error in create symlink block in \n");
27         goto bad;
28     }
29     iunlock(dp);
30 }
31 ilock(dp);
32 // dp holds the new path inode
33 if(writei(dp,0,(uint64)old,dp->size,MAXPATH) != MAXPATH){
34     printf("write error in symlink\n");
35     iunlock(dp);
36     goto bad;
37 }
38 //printf("%d -> %d\n", dp->inum, );
39 dp->type=T_SYMLINK;
40 iunlockput(dp);
41 end_op();
42
43 return 0;
44 //must release all inode lock before going to bad
45 bad:
46     end_op();
47     return -1;
48
49 }
50

```

2. recursively follow the link to find to node

```

1  struct inode * fetchSym(char *path, int depth){
2      if (depth > 8){
3          return 0;
4      }
5      struct inode *ip;
6      if((ip = namei(path)) == 0){
7          return 0;
8      }
9      ilock(ip);
10     if(ip->type==T_SYMLINK ){
11         char next[MAXPATH];
12         if(readi(ip,0,(uint64)next,ip->size-MAXPATH ,MAXPATH) == 0 ) {
13             iunlock(ip);
14             return 0;
15         }
16         iunlock(ip);
17         return fetchSym(next, depth+1 );
18     }
19     iunlock(ip);
20     return ip;
21 }

```

## Result

Pass all test

```

make[1]: Leaving directory '/home/vielo/code/xv6lab'
== Test running bigfile ==
$ make qemu-gdb
running bigfile: OK (120.7s)
== Test running symlinktest ==
$ make qemu-gdb
(1.0s)
== Test  symlinktest: symlinks ==
symlinktest: symlinks: OK
== Test  symlinktest: concurrent symlinks ==
symlinktest: concurrent symlinks: OK
== Test usertests ==
$ make qemu-gdb
usertests: OK (221.5s)
== Test time ==
time: OK
Score: 100/100

```

