

Mmap

Design

You should implement enough `mmap` and `munmap` functionality to make the `mmaptest` test program work. If `mmaptest` doesn't use a `mmap` feature, you don't need to implement that feature.

What is mmap?

The `mmap` and `munmap` system calls allow UNIX programs to exert detailed control over their address spaces.

In this lab, my job is to mapping files in fs to memory space and sync these 2 parts.

How to mmap ?

1. maintain all vma:

```
1  struct vma{
2      struct spinlock lock;
3      int pid;
4      uint64 rootva;
5      uint64 start;
6      uint64 end;
7      //only when changing length the lock will be held
8      int length;
9      int prot;
10     int flags;
11     struct file *f;
12     struct vma *next;
13 };
14 struct vma vmaPool[MAXMMAP];
15
```

2. allocate a new vma

```
1  //alloc a vma
2  struct vma *v = vmaAlloc();
3  //printf("alloc a vma ok\n");
4  if (v==0){
5      printf("no vma block remain\n");
6      return -1;
7  }
```

```

8    //printf("lock dddd \n");
9    v->pid = p->pid;
10   v->rootva = start;
11   v->start = start;
12   v->end = PGROUNDUP(start + size);
13   acquire(&v->lock);
14   v->length = v->end - v->start;
15   release(&v->lock);
16   v->prot = pteFlag;
17   v->flags = flags;
18   v->f = f;
19   v->next = 0;
20
21   //add new vma to pcb
22   struct vma * vp = p->vp;
23   if (vp==0){
24       p->vp = v;
25       //printf("l1 vma %d, l2 vma %d\n",p->vp,p->vp->next);
26   }else{
27       while (vp->next!=0){
28           //printf("go down\n");
29           vp = vp->next;
30       }
31       vp->next = v;
32   }

```

3. allocate free space in process's pagetable

```

1    p->sz = v->end;
2    printf("map from %p to %p\n",v->start,v->end);

```

Handle page fault

If process wanna access to this vma page, it will throw a page fault. my job is to read the data from file, put it to newly-allocated physiscal page for that va, and mapping them.

```

1    int mmapHandler(uint64 va, uint64 cause){
2        struct proc *p = myproc();
3        struct vma* v = p->vp;
4        if (va >= MAXVA || va ==0){
5            return -1;
6        }
7        while(v != 0){
8            if(va >= v->start && va < v->end){
9                goto found ;
10           }
11           v = v->next;

```

```

12     }
13     printf("addr not in vma\n");
14     return -1;
15 found:
16
17     if(cause == 13 && (v->prot & PTE_R)==0){
18         //read unreadable vma
19         return -1;
20     }
21     if(cause == 15 && (v->prot & PTE_W)==0){
22         //write unwriteable vma
23         return -1;
24     }
25
26     char * mem = kalloc();
27     if (mem==0){
28         printf("no physical memory for vma mapping\n");
29         return -1;
30     }
31     memset(mem,0,PGSIZE);
32     int pteFlags = v->prot ;
33     if (cause == 15){
34         pteFlags |= PTE_D;
35     }
36     if( mappages(p->pagetable,PGROUNDOWN(va),PGSIZE,(uint64)mem,pteFlags) !=0 ) {
37         kfree(mem);
38         printf("mapping error in mmap trap\n");
39         return -1;
40     }
41     struct file *fl = v->f;
42     ilock(fl->ip);
43     readi(fl->ip, 0, (uint64)mem, PGROUNDOWN(va) - v->start, PGSIZE);
44     iunlock(fl->ip);
45     printf("load va %p 's data in memory from fs\n",va);
46
47     return 0;
48 }

```

How to unmmmap?

There are different free method:

1. free the front part of vma address space
2. free the back part of vma address space
3. free a whole vma address space

1. modify vma:

```

1     if(va == v->start){

```

```

2     if (vaEnd == v->end){
3         //free all mapping
4         //free this vma to pool
5         //fclose(v->f);
6         v->start = 0;
7         v->end = 0;
8         if(pv==0){
9             //only 1 vma
10            p->vp = 0;
11        }else{
12            pv->next = v->next;
13        }
14    }else {
15        //free a part of the front
16        v->start = vaEnd;
17    }
18 }else{
19     //free a part of the end
20     v->end = va;
21 }

```

2. sync unmap memory back to file in fs

```

1 //write back [va,vaEnd] to fs
2 pte_t * pte;
3 if( (v->prot & PTE_W) && (v->flags & MAP_SHARED) ){
4     for (uint64 addr = va; addr < vaEnd; addr+= PGSIZE){
5         pte = walk(p->pagetable,addr,0);
6         if (pte !=0 && (*pte & PTE_D)){
7             writebackPage(v->f,addr,PGSIZE,addr - v->rootva);
8         }
9     }
10 }

```

3. free the cache memory in process's pagetable

```

1 //free unmap pages
2 uvmunmap(p->pagetable,va,(vaEnd -va)/PGSIZE,1);
3 if(v->start == v->end){
4     fclose(v->f);
5 }
6 acquire(&v->lock);
7 v->length = v->end -v->start;
8 release(&v->lock);
9 printf("after unmap %p,%p\n",v->start,v->end);
10 return 0;

```

Result

pass all tests

```
nel.sym
make[1]: Leaving directory '/home/vielo/cod
== Test running mmaptest ==
$ make qemu-gdb
(5.5s)
== Test    mmaptest: mmap f ==
    mmaptest: mmap f: OK
== Test    mmaptest: mmap private ==
    mmaptest: mmap private: OK
== Test    mmaptest: mmap read-only ==
    mmaptest: mmap read-only: OK
== Test    mmaptest: mmap read/write ==
    mmaptest: mmap read/write: OK
== Test    mmaptest: mmap dirty ==
    mmaptest: mmap dirty: OK
== Test    mmaptest: not-mapped unmap ==
    mmaptest: not-mapped unmap: OK
== Test    mmaptest: two files ==
    mmaptest: two files: OK
== Test    mmaptest: fork_test ==
    mmaptest: fork_test: OK
== Test usertests ==
$ make qemu-gdb
usertests: OK (103.2s)
== Test time ==
time: OK
Score: 140/140
(base)
```